



Introduction To Pig

Alan Gates

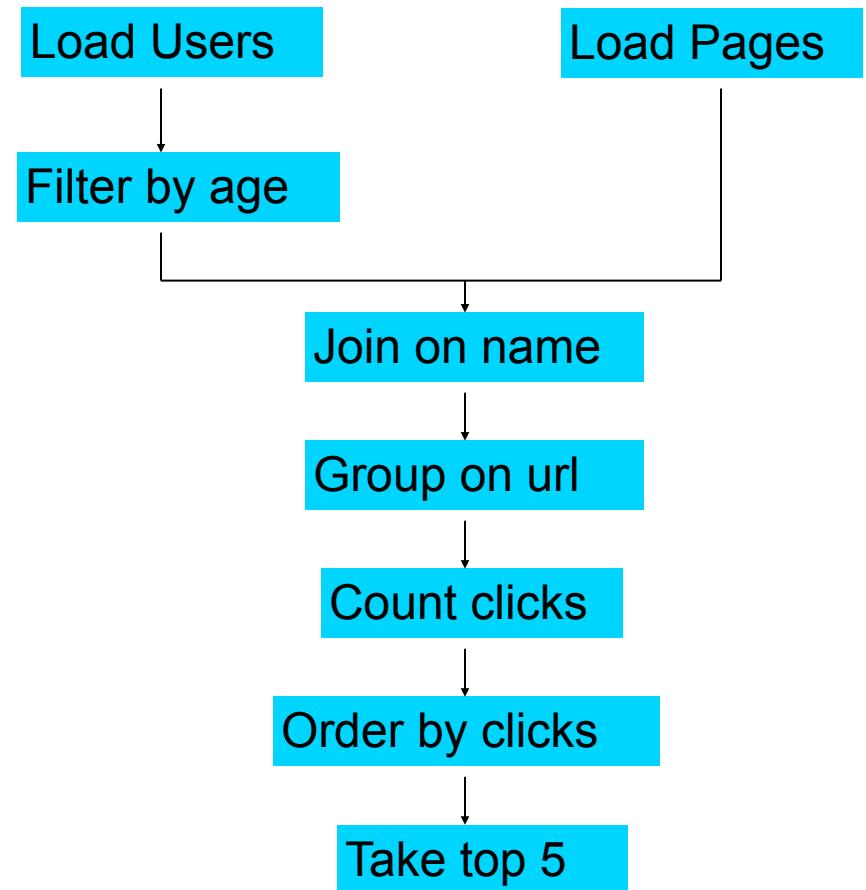
Yahoo!

About Me

- Pig Committer
- Member of Hadoop PMC
- Architect for Pig development group at Yahoo!

Why is Pig Interesting?

- Data
 - User records
 - Pages served
- Question: the 5 pages most visited by users aged 18 - 25.



..



In Map Reduce

```

import java.io.IOException;
import java.util.ArrayList;
import java.util.Iterator;
import java.util.List;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;
import org.apache.hadoop.io.WritableComparable;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapred.TextInputFormat;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.RedException;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.SequenceFileInputFormat;
import org.apache.hadoop.mapred.SequenceFileOutputFormat;
import org.apache.hadoop.mapred.TextInputFormat;
import org.apache.hadoop.mapred.jobcontrol.Job;
import org.apache.hadoop.mapred.jobcontrol.JobControl;
import org.apache.hadoop.mapred.lib.IdentityMapper;
public class MRExample {
    public static class LoadPages extends MapReduceBase
        implements Mapper<LongWritable, Text, Text, Text> {
        public void map(LongWritable k, Text val,
            OutputCollector<Text, Text> oc,
            Reporter reporter) throws IOException {
            // Pull the key out
            String line = val.toString();
            int firstComma = line.indexOf(',');
            String value = line.substring(0, firstComma);
            Text outKey = new Text(key);
            // Prepend an index to the value so we know which file
            // it came from.
            Text outVal = new Text("1 " + value);
            oc.collect(outKey, outVal);
        }
    }
    public static class LoadAndFilterUsers extends MapReduceBase
        implements Mapper<LongWritable, Text, Text, Text> {
        public void map(LongWritable k, Text val,
            OutputCollector<Text, Text> oc,
            Reporter reporter) throws IOException {
            // Pull the key out
            String line = val.toString();
            int firstComma = line.indexOf(',');
            String value = line.substring(0, firstComma);
            int age = Integer.parseInt(value);
            if (age < 18 || age > 25) return;
            String key = line.substring(firstComma + 1);
            // Add an index to the value so we know which file
            // it came from.
            Text outVal = new Text("2 " + value);
            oc.collect(outKey, outVal);
        }
    }
    public static class Join extends MapReduceBase
        implements Reducer<Text, Text, Text, Text> {
        public void reduce(Text key,
            Iterator<Text> iter,
            OutputCollector<Text, Text> oc,
            Reporter reporter) throws IOException {
            // For each value, figure out which file it's from and
            store it
            // accordingly.
            List<String> first = new ArrayList<String>();
            List<String> second = new ArrayList<String>();
            while (iter.hasNext()) {
                Text t = iter.next();
                String value = t.toString();
                if (value.charAt(0) == '1')
                    first.add(value.substring(1));
                else second.add(value.substring(1));
            }
        }
    }
    public static class LoadJoined extends MapReduceBase
        implements Mapper<Text, Text, Text, LongWritable> {
        public void map(
            Text k,
            Text val,
            OutputCollector<Text, LongWritable> oc,
            Reporter reporter) throws IOException {
            // Find the url
            String line = val.toString();
            int firstComma = line.indexOf(',');
            int secondComma = line.indexOf(',', firstComma);
            String url = line.substring(firstComma, secondComma);
            // Drop the rest of the record, I don't need it anymore,
            // just pass a 1 for the combiner/reducer to sum instead.
            Text outKey = new Text(key);
            oc.collect(outKey, new LongWritable(1L));
        }
    }
    public static class ReduceUrls extends MapReduceBase
        implements Reducer<Text, LongWritable, WritableComparable,
        Writable> {
        public void reduce(
            Text key,
            Iterator<LongWritable> iter,
            OutputCollector<WritableComparable, Writable> oc,
            Reporter reporter) throws IOException {
            // Add up all the values we see
            long sum = 0;
            while (iter.hasNext()) {
                sum += iter.next().get();
                reporter.setStatus("OK");
            }
            oc.collect(key, new LongWritable(sum));
        }
    }
    public static class LoadClicks extends MapReduceBase
        implements Mapper<WritableComparable, Writable, LongWritable,
        Text> {
        public void map(
            WritableComparable key,
            Writable val,
            OutputCollector<LongWritable, Text> oc,
            Reporter reporter) throws IOException {
            oc.collect((LongWritable)val, (Text)key);
        }
    }
    public static class LimitClicks extends MapReduceBase
        implements Reducer<LongWritable, Text, LongWritable, Text> {
        int count = 0;
        public void reduce(
            LongWritable key,
            Iterator<Text> iter,
            OutputCollector<LongWritable, Text> oc,
            Reporter reporter) throws IOException {
            // Only output the first 100 records
            while (count < 100 && iter.hasNext()) {
                oc.collect(key, iter.next());
                count++;
            }
        }
    }
    public static void main(String[] args) throws IOException {
        JobConf lp = new JobConf(MRExample.class);
        lp.setJobName("Load Pages");
        lp.setInputFormat(TextInputFormat.class);
        lp.setOutputKeyClass(Text.class);
        lp.setOutputValueClass(Text.class);
        lp.setMapperClass(LoadPages.class);
        FileInputFormat.addInputPath(lp, new
        Path("/user/gates/pages"));
        FileOutputFormat.setOutputPath(lp,
            new Path("/user/gates/tmp/indexed_pages"));
        lp.setNumReduceTasks(0);
        Job loadPages = new Job(lp);
        JobConf lfu = new JobConf(MREExample.class);
        lfu.setJobName("Load and Filter Users");
        lfu.setInputFormat(TextInputFormat.class);
        lfu.setOutputKeyClass(Text.class);
        lfu.setOutputValueClass(Text.class);
        lfu.setMapperClass(LoadAndFilterUsers.class);
        FileInputFormat.addInputPath(lfu, new
        Path("/user/gates/users"));
        FileOutputFormat.setOutputPath(lfu,
            new Path("/user/gates/tmp/fILTERED_USERS"));
        lfu.setNumReduceTasks(0);
        lfu.setJobName("Load and Filter Users");
        JobConf join = new JobConf(MREExample.class);
        join.setJobName("Join Users and Pages");
        join.setInputFormat(TextInputFormat.class);
        join.setOutputKeyClass(Text.class);
        join.setOutputValueClass(Text.class);
        join.setMapperClass(Join.class);
        join.setReducerClass(Join.class);
        FileInputFormat.addInputPath(join, new
        Path("/user/gates/tmp/indexed_PAGES"));
        FileInputFormat.addInputPath(join, new
        Path("/user/gates/tmp/filterd_USERS"));
        FileOutputFormat.setOutputPath(join, new
        Path("/user/gates/tmp/joined"));
        join.setNumReduceTasks(50);
        join.setJobName("join");
        joinJob.addDependingJob(loadPages);
        joinJob.addDependingJob(loadUsers);
        JobConf groupJob = new JobConf(MREExample.class);
        groupJob.setJobName("Group URLs");
        groupJob.setInputFormat(KeyValueTextInputFormat.class);
        group.setOutputKeyClass(Text.class);
        group.setOutputValueClass(LongWritable.class);
        group.setMapperClass(LoadJoined.class);
        group.setCombinerClass(LimitClicks.class);
        group.setReducerClass(ReduceUrls.class);
        FileInputFormat.addInputPath(group, new
        Path("/user/gates/tmp/joined"));
        FileOutputFormat.setOutputPath(group, new
        Path("/user/gates/tmp/grouped"));
        group.setNumReduceTasks(50);
        Job groupJob = new Job(group);
        groupJob.addDependingJob(joinJob);
        JobConf top100 = new JobConf(MREExample.class);
        top100.setJobName("Top 100 sites");
        top100.setInputFormat(TextInputFormat.class);
        top100.setOutputKeyClass(LongWritable.class);
        top100.setOutputValueClass(Text.class);
        top100.setMapperClass(LoadClicks.class);
        top100.setCombinerClass(LimitClicks.class);
        top100.setReducerClass(Summarize.class);
        FileInputFormat.addInputPath(top100, new
        Path("/user/gates/tmp/grouped"));
        FileOutputFormat.setOutputPath(top100, new
        Path("/user/gates/top100sitesforusers18to25"));
        top100.setNumReduceTasks(1);
        Job limit = new Job(top100);
        limit.addDependingJob(groupJob);
        JobControl jc = new JobControl("Find top 100 sites for users
        18 to 25");
        jc.addJob(loadPages);
        jc.addJob(loadUsers);
        jc.addJob(joinJob);
        jc.addJob(groupJob);
        jc.addJob(limit);
        jc.run();
    }
}

```



In Pig Latin

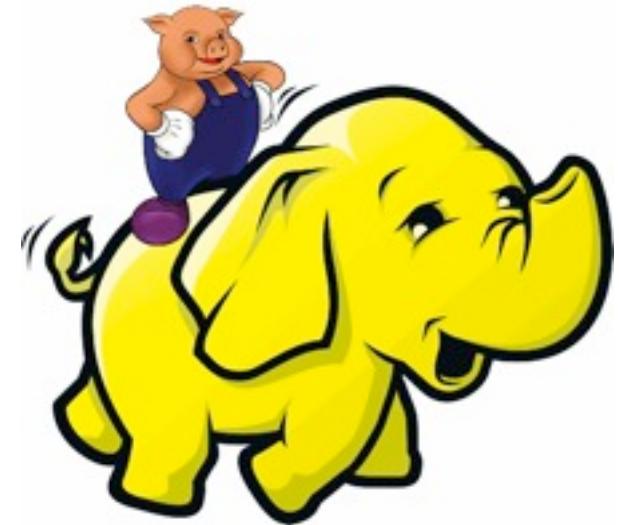
```
Users = load 'users' as (name, age);  
Fltrd = filter Users by  
    age >= 18 and age <= 25;  
Pages = load 'pages' as (user, url);  
Jnd = join Fltrd by name, Pages by user;  
Grpd = group Jnd by url;  
Smmd = foreach Grpd generate group,  
    COUNT(Jnd) as clicks;  
Srted = order Smmd by clicks desc;  
Top5 = limit Srted 5;  
store Top5 into 'top5sites';
```

What is Pig?

Pig Latin, a high level data processing language.



An engine that executes Pig Latin locally or on a Hadoop cluster.



A subproject of
Apache Hadoop



Comparing Pig and Map Reduce

- Pig
 - radically reduces development time
 - includes standard relational operations (join, sort, etc.) so you don't have to write them yourself
 - copes with Hadoop version and API changes
- While preserving the power of Hadoop
 - user code can be plugged in at almost any point
 - works with structured, semi-structured, and unstructured data
 - no requirement up front to define data schema
 - explicit data flow control

Why not SQL?

- SQL is spoken almost everywhere
- But
 - SQL requires schema definition up front
 - SQL lets you say what you want, but not how you want it done
 - SQL works on tables, but HDFS has files
 - Writing long involved data processes in SQL is painful
(subqueries anyone?)
- Pig Latin seeks to be the native language for data processing in Hadoop

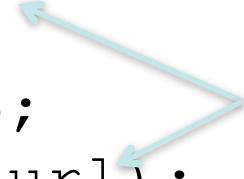
Detailed Example

```
Users = load 'users' as (name, age);  
Fltrd = filter Users by  
    age >= 18 and age <= 25;  
Pages = load 'pages' as (user, url);  
Jnd = join Fltrd by name, Pages by user;  
Grpd = group Jnd by url;  
Smmd = foreach Grpd generate group,  
    COUNT(Jnd) as clicks;  
Srted = order Smmd by clicks desc;  
Top5 = limit Srted 5;  
store Top5 into 'top5sites';
```



Detailed Example

```
Users = load 'users' as (name, age);  
Fltrd = filter Users by  
    age >= 18 and age <= 25;  
Pages = load 'pages' as (user, url);  
Jnd = join Fltrd by name, Pages by user;  
Grpd = group Jnd by url;  
Smmd = foreach Grpd generate group,  
    COUNT(Jnd) as clicks;  
Srted = order Smmd by clicks desc;  
Top5 = limit Srted 5;  
store Top5 into 'top5sites';
```



Schema defined at runtime

Detailed Example

```
Users = load 'users' as (name, age);  
Fltrd = filter Users by  
    age >= 18 and age <= 25;  
Pages = load 'pages' as (user, url);  
Jnd = join Fltrd by name, Pages by user;  
Grpd = group Jnd by url;  
Smmd = foreach Grpd generate group,  
    COUNT(Jnd) as clicks;  
Srted = order Smmd by clicks desc;  
Top5 = limit Srted 5;  
store Top5 into 'top5sites';
```



Detailed Example

```
Users = load 'users' as (name, age);
Fltrd = filter Users by
          age >= 18 and age <= 25;           Dynamic
                                                typing
Pages = load 'pages' as (user, url);
Jnd = join Fltrd by name, Pages by user;
Grpd = group Jnd by url;
Smmd = foreach Grpd generate group,
        COUNT(Jnd) as clicks;
Srted = order Smmd by clicks desc;
Top5 = limit Srted 5;
store Top5 into 'top5sites';
```



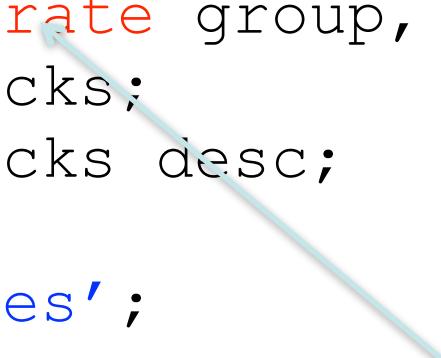
Detailed Example

```
Users = load 'users' as (name, age);  
Fltrd = filter Users by  
    age >= 18 and age <= 25;  
Pages = load 'pages' as (user, url);  
Jnd = join Fltrd by name, Pages by user;  
Grpd = group Jnd by url;  
Smmd = foreach Grpd generate group,  
    COUNT(Jnd) as clicks;  
Srted = order Smmd by clicks desc;  
Top5 = limit Srted 5;  
store Top5 into 'top5sites';
```



Detailed Example

```
Users = load 'users' as (name, age);  
Fltrd = filter Users by  
    age >= 18 and age <= 25;  
Pages = load 'pages' as (user, url);  
Jnd = join Fltrd by name, Pages by user;  
Grpd = group Jnd by url;  
Smmd = foreach Grpd generate group,  
    COUNT(Jnd) as clicks;  
Srted = order Smmd by clicks desc;  
Top5 = limit Srted 5;  
store Top5 into 'top5sites';
```

yahoo.com, 
$$\begin{cases} (\text{joe}, 19), \\ (\text{bob}, 21), \\ (\text{mary}, 22) \end{cases}$$



Plugging in Your Code

```
A = load 'webpages' using JsonLoader();  
B = foreach B generate Canonical(url), content;  
C = filter B by NotABot(url);  
C = group B by url;  
D = foreach C generate group, PageRank(D1);  
store D into 'rankedpages' using AvroStorage();
```

Plugging in Your Code

User can define how data is parsed

```
A = load 'webpages' using JsonLoader();  
B = foreach B generate Canonical(url), content;  
C = filter B by NotABot(url);  
C = group B by url;  
D = foreach C generate group, PageRank(D1);  
store D into 'rankedpages' using AvroStorage();
```



Plugging in Your Code

```
A = load 'webpages' using JsonLoader();  
B = foreach B generate Canonical(url), content;  
C = filter B by NotABot(url);  
C = group B by url;  
D = foreach C generate group, PageRank(D1);  
store D into 'rankedpages' using AvroStorage();
```

Plugging in Your Code

Loader can provide schema

```
A = load 'webpages' using JsonLoader();  
B = foreach B generate Canonical(url), content;  
C = filter B by NotABot(url);  
C = group B by url;  
D = foreach C generate group, PageRank(D1);  
store D into 'rankedpages' using AvroStorage();
```



Plugging in Your Code

```
A = load 'webpages' using JsonLoader();  
B = foreach B generate Canonical(url), content;  
C = filter B by NotABot(url);  
C = group B by url;  
D = foreach C generate group, PageRank(D1);  
store D into 'rankedpages' using AvroStorage();
```

Plugging in Your Code

Column transformation
function

```
A = load 'webpages' using JsonLoader();  
B = foreach B generate Canonical(url), content;  
C = filter B by NotABot(url);  
C = group B by url;                                ← Filter function  
D = foreach C generate group, PageRank(D1);  
store D into 'rankedpages' using AvroStorage();
```

Aggregate



Plugging in Your Code

```
A = load 'webpages' using JsonLoader();  
B = foreach B generate Canonical(url), content;  
C = filter B by NotABot(url);  
C = group B by url;  
D = foreach C generate group, PageRank(D1);  
store D into 'rankedpages' using AvroStorage();
```

Plugging in Your Code

```
A = load 'webpages' using JsonLoader();  
B = foreach B generate Canonical(url), content;  
C = filter B by NotABot(url);  
C = group B by url;  
D = foreach C generate group, PageRank(D1);  
store D into 'rankedpages' using AvroStorage();
```



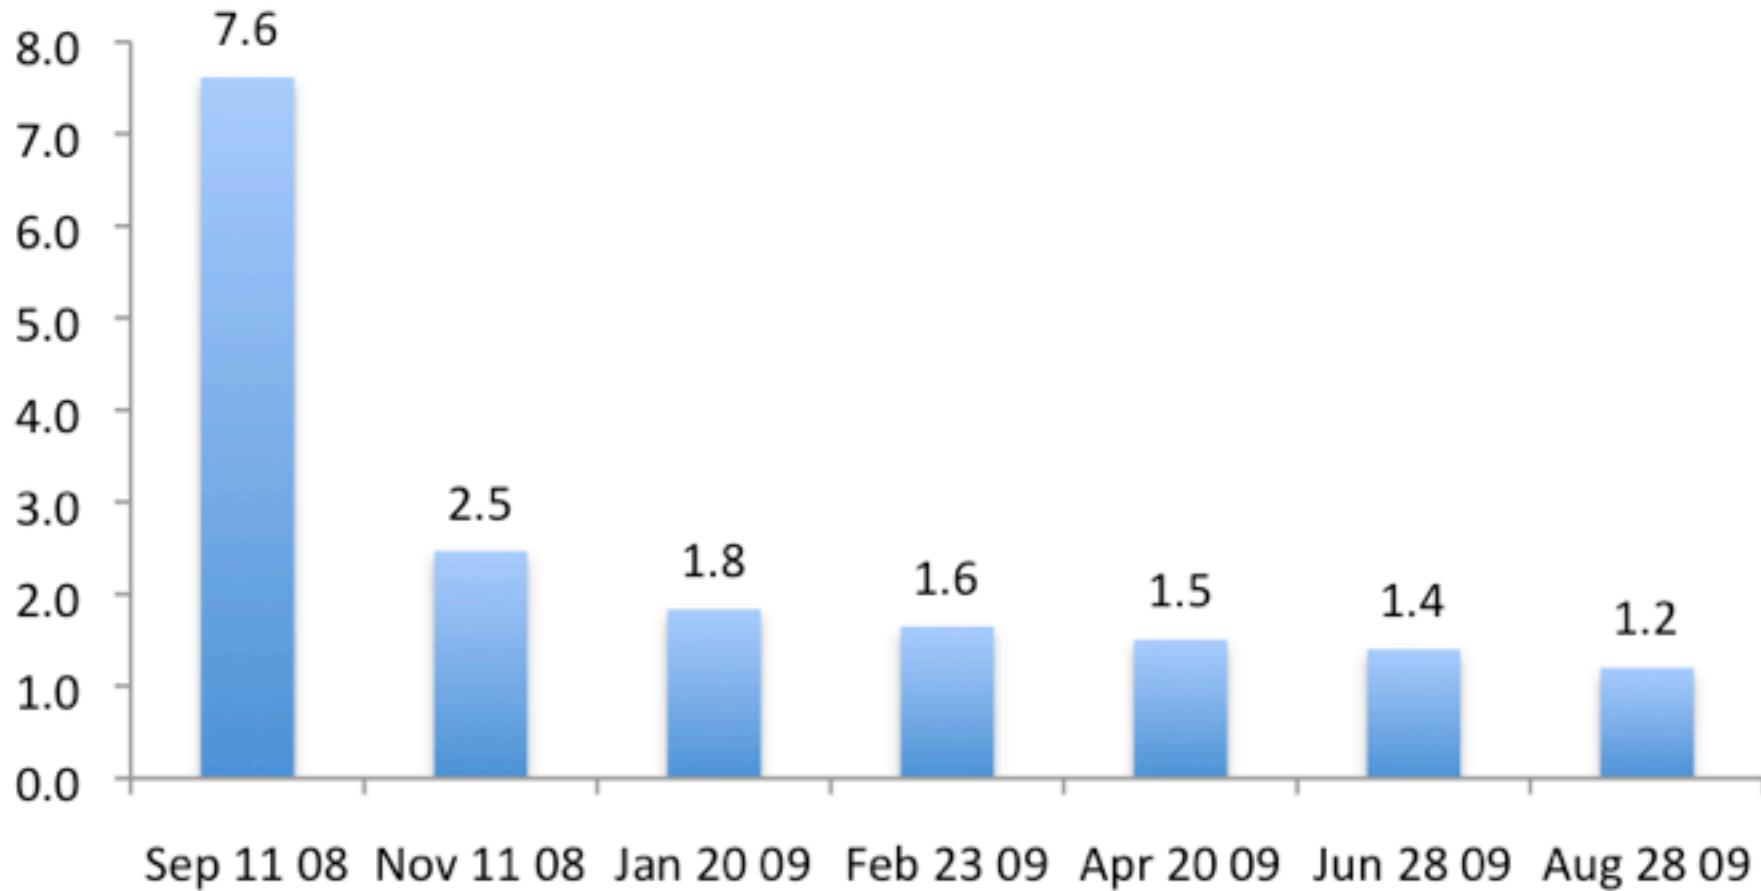
User can define how data is stored

Branching Data Flows

```
A = load 'users' as (name, age, gender,  
    city, state);  
B = filter A by name is not null;  
C = group B by age, gender;  
D = foreach C generate group, COUNT(B);  
store D into 'by_demographic';  
gamma = group B by state;  
delta = foreach gamma generate group, COUNT(B);  
store delta into 'by_state';
```

Performance

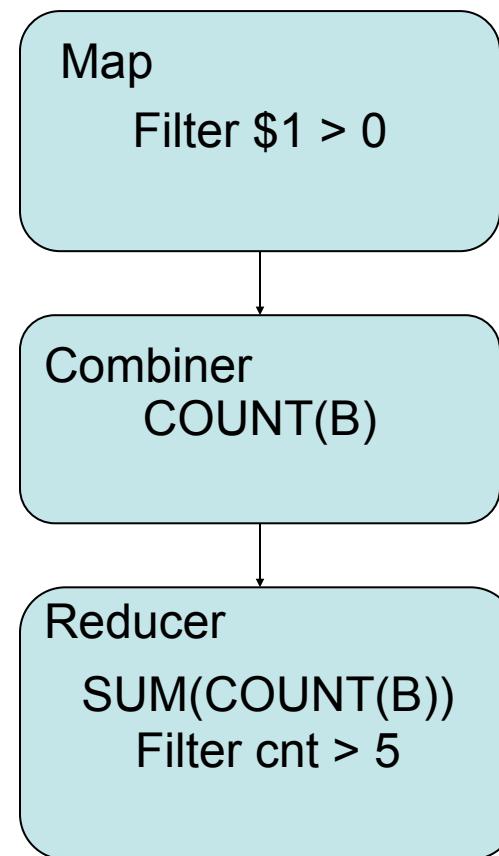
Pig Performance vs Map-Reduce



How it Works

Pig Latin script is translated to a set of operators which are placed in one or more MR jobs and executed.

```
A = load 'myfile';
B = filter A by $1 > 0;
C = group B by $0;
D = foreach C generate
    group, COUNT(B) as cnt;
E = filter D by cnt > 5;
dump E;
```



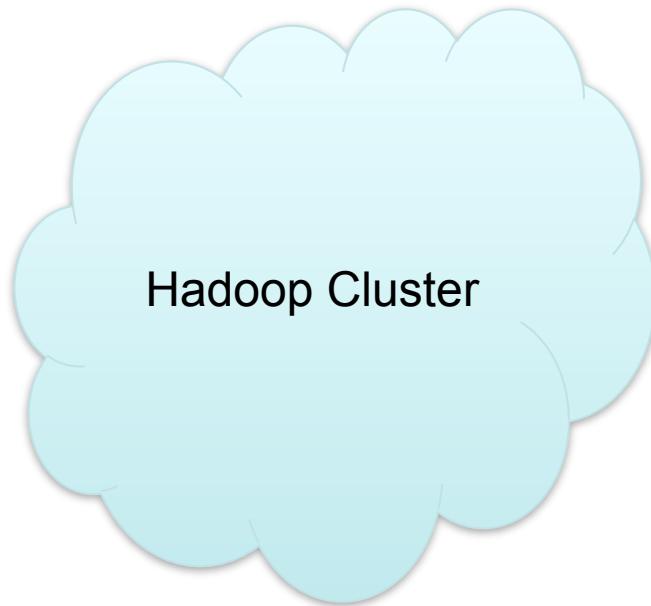
Components

Pig resides on user machine



User machine

Job executes on cluster



No need to install anything extra on your Hadoop cluster.

What Users Are Doing With Pig

- Inside Yahoo (based on user interviews)
 - 60% of ad hoc and 40% of production MR jobs
 - Production
 - Examples: search infrastructure, ad relevance
 - Attraction: fast development, extensibility via custom code, protection against Hadoop changes, debugability
 - Ad hoc
 - Examples: user intent analysis
 - Attraction: easy to learn, compact readable code, fast iteration when trying new algorithms, easy for collaboration

What Users Are Doing With Pig

- Outside Yahoo (based on mailing list responses)
 - Processing search engine query logs

“Pig programs are easier to maintain, and less error-prone than native java programs. It is an excellent piece of work.”
 - Image recommendations

“I am using it as a rapid-prototyping language to test some algorithms on huge amounts of data.”
 - Adsorption Algorithm (video recommendations)
 - Hoffman’s PLSI implementation

“The E/M login was implemented in pig in 30-35 lines of pig-latin statements. Took a lot less compared to what it took in implementing the algorithm in mapreduce java. Exactly that’s the reason I wanted to try it out in Pig. It took ~ 3-4 days for me to write it, starting from learning pig.”

Up and Coming Features

- UDFs in scripting languages that compile to Java byte code (JRuby, Jython, Groovy)
- Making Pig bilingual: SQL
- Performance: faster, faster, faster
- Better integration with other Hadoop subprojects

..



Getting Your Own Pig

- You can download Pig from <http://hadoop.apache.org/pig/>
- Pig is included as part of the Cloudera distribution <http://www.cloudera.com/hadoop>
- Amazon Elastic Map Reduce service allows you to use Hadoop with no hardware investment, and supports Pig:
<http://aws.amazon.com/elasticmapreduce/>

Next Steps

- Read the online documentation: <http://hadoop.apache.org/pig/>
- On line tutorials
 - From Yahoo, <http://developer.yahoo.com/hadoop/tutorial/>
 - From Cloudera, <http://www.cloudera.com/hadoop-training>
- A couple of Hadoop books available that include chapters on Pig, search at your favorite bookstore
- Join the mailing lists:
 - pig-user@hadoop.apache.org for user questions
 - pig-dev@hadoop.apache.com for developer issues

Questions

